

Minimizing MSE for Linear Models

****BACK TO. 2-5 NOTES PAGE 2**** Complexity of solving equations or finding matrix inverse is $O(Nd^2 + d^3)$

Minimizing MSE for Non-Linear Equations

$MSE(\theta)$ is concave so you can use gradient descent:

$$\theta^{new} = \theta^{current} - stepSize * gradient(MSE)$$

Stochastic Gradient Descent can also be used, faster but noisier

Probabilistic Interpretation of Regression

$p(y|x)$: for fixed x , there is variation in y 2 types of variation: - Measurement noise - Unobserved Variables
2 sources of variability: $p(y|x)$: variability in y given x $p(x)$: distribution of input data in the space We have a joint distribution: $p(x, y) = p(y|x)p(x)$ and we learn $p(y|x)$

Modeling Framework

$y_x = E[y|x] + e$ where

y_x : what we observe

$E[y|x]$: what we try to learn with $f(x, \theta)$

e : unpredictable error term

Simple Model:

$p(y|x) = N(f(x, \theta), \sigma^2)$ where $f(x, \theta) = \theta^T x$

Conditional Likelihood for Regression

$$L(\theta) = \prod p(y_i|x_i, \theta)$$

As an example, we can use a Gaussian model for $p(y|x, \theta)$.

This will give us

$$p(y|x, \theta) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2\sigma^2}(y - f(x, \theta))^2\right)$$

This equation comes from the fact that $E[y|x] = f(x, \theta)$. After doing some algebra

$$\log L(\theta) \propto -MSE(\theta)$$

Thus maximizing log likelihood is the same as minimizing MSE. This gives us a useful framework to go beyond.

Bayesian View of Regression

Posterior Density on θ is:

$$p(\theta|D_x, D_y) \propto p(D_x, D_y|\theta)p(\theta) = p(D_y|D_x, \theta)p(D_x|\theta)p(\theta)$$

$$p(\theta|D_x, D_y) \propto p(D_y|D_x, \theta)p(\theta)$$

This is because we do not model $p(D_x|\theta)$

****TODO: BACK TO GAUSSIAN ERROR MODEL IN 2-19 NOTES****

Properties of Minimizing MSE

Because of this fact:

$$MSE(\theta) = \int \int (y - f(x; \theta))^2 p(x, y) dx dy$$

It eventually (****TODO: make sure you can derive this. 2-19 notes****) holds that the optimal value is when $f(x; \theta) = E[y_x]$.

We are limited because

Bias: $f(x, \theta)$ might not be able to exactly approximate $E[y_x]$

Variance: Even if so, only have finite data to learn $E[y_x]$

Tradeoff exists between complex model with low bias but high variance and simple model with high bias and low variance

Bias-Variance Tradeoff

****TODO: DERIVE THIS FROM 2-19 NOTES**** Eventually we have $MSE_x = \sigma_{yx}^2 + Bias^2 + Variance$

This leads to the fundamental bias-variance tradeoff

- simple models with few parameters have high bias, but low variance
- complex models with many parameters have low bias, but high variance

Logistic Regression Classifier

Regression but $y_i \in 0, 1$ and $p(y|x)$ is the logistic function

****MAKE SURE TO UNDERSTAND WHERE LOGISTIC FUNCTION COMES FROM****

****MAKE SURE TO REMEMBER FORMULAS AND IDEA FROM 2-24 NOTES****

****BACK TO MULTI-CLASS LOGISTIC REGRESSION AND CONNECTION TO FEED-FORWARD NEURAL NETWORKS ON 2-26 NOTES****

Generative Approach to Classification

We model $p(x|c)$ instead of $p(c|x)$. Models are generative when we model the distribution of both x's and c's.

$$L(\theta) = \prod p(x_i|c_i, \theta)p(c_i)$$

Key Points:

- learn how x_i values are distributed for each class, so θ_k is set of parameters for class k model.
- learn $p(c=k)$
- possibly decomposes into k optimization problems
- it is optimal if distributional assumptions are correct
- predict using bayes rule

$$p(c = k|x, \theta) \propto p(x|c = k, \theta_k)p(c = k)$$

or we could be Bayesian and average over the θ values.

Weaknesses of Gaussian model for each class:

- sensitive to Gaussian assumption
- scales poorly as d increases. for high dimensions, we can assume covariance matrices are diagonal.

Naive Bayes Model:

If x_i are binary vectors, we can use a Naive Bayes model for each class. Parameters are $\theta_k = \{\theta_{k1}, \dots, \theta_{kd}\}$ where θ_{kj} is Bernoulli probability that $x_{ij} = 1$

FIRST ORDER MARKOV MODEL EXAMPLE WILL NOT BE ON THE FINAL*

Discriminant Functions

To make a decision about mostly class, we can compute $\operatorname{argmax}_k p(c = k|x)$

Using Bayes rule, this is $\operatorname{argmax}_k p(x|c = k)p(c = k)$

This is equal to $\operatorname{argmax}_k \log(p(x|c = k)) + \log(p(c = k))$

All of these can be used as discriminant functions $g_k(x)$

For 2 class case, decision boundary is when $g(x) = g_1(x) - g_2(x) = 0$

****BACK TO MULTIVARIATE GAUSSIAN DISCRIMINANT FUNCTION FROM 2-26****

Finite Mixture Models Definition

Definition is

$$p(x) = \sum p(x, z) = \sum p_k(x|\theta_k, z = k)p(z = k)$$

where $\sum p(z = k) = 1$

The z variable is hidden and the densities can be a mixture

Learning Mixture Models

Given data $D = x_i$ and form of each component $p_k(x|\theta_k)$.

Pseudo-code:

for i = 1:N

- sample component $p(z = k)$ for i-th data point, call it k^*
- sample x_i from k^* . $p_k(x|\theta_k, z = k^*)$

Let α_k be the component weight, then log likelihood is

$$l(\theta) = \sum \log(p(x_i|\theta)) = \sum \log \sum p_k(x_i|\theta_k)\alpha_k$$

BACK TO EM ALGORITHM PROCEDURE FROM 3-5 NOTES*

Notes about EM

- EM converges to a local minimum of the likelihood or log-likelihood. We are essentially doing a search in θ space.

- similar to gradient ascent but more popular because step size is chosen automatically

EM respects parameter constraints

The rate at which EM converges is a function of how much missing info is in the problem. Two Gaussians close together has a lot of missing info, but if they are far apart, the uncertainty is lower.

EM can be mixed with other methods

Variations of EM

Online EM: in the E and M step only use a subset of points. Useful when N is large. After many noisy steps, it eventually converges

Generalized EM: *BACK TO. IN 3-5 NOTES*

Approximate E steps: in some models, computing the likelihood in the E step so we can replace the E step with an approximation, such as Monte Carlo sampling

Semi-supervised learning: some x_i have labels and some do not. This is handled nicely in EM since we can fix our weights for the x_i that have values.

Singular Solutions in EM

To avoid it, common solution is to set lower limit on σ_k^2 .

Another method is to put priors $p(\theta)$ on parameters and extend EM to compute MAP estimates instead of ML in the M step. Gives us a weighted average of our usual estimate

$$\Sigma_{MAP,k} = \alpha \Sigma_k^{ML} + (1 - \alpha) \Sigma_0$$

The prior Σ_0 is usually chosen to be diagonal.

Kernal Density Estimation

*GO BACK TO. IN 3-10 NOTES**

Hidden Markov Models

At each time t , z_t produces a new data point x_t and then transitions to a new state z_{t+1}

Two assumptions:

observations x_t are CI of all other variables given z_t so observation at time t depends only on the current state

the z_t values form a first order Markov Chain.

Equation from graphical model is the following

$$p(x_{1..T}, z_{1..T}) = \prod p(x_t | z_t) p(z_t | z_{t-1})$$

Two sets of parameters: **BACK TO. IN 3-10 NOTES**

Computing Likelihood of HMMs

BACK TO EFFICIENT COMPUTATION OF HMM LIKELIHOOD IN 3-10 NOTES AND SOME IN 3-12 NOTES

FINAL MATERIAL ENDS HERE. EM FOR HMM WILL NOT BE ON FINAL